Today’s material available at:

https://github.com/juliennelachance/ai4all_clustering

Clustering

Princeton AI4ALL

By Julie LaChance

Icebreaker

- Everyone receives a slip of paper
- Your task:

  Using the information about your person, form dog-walking groups
- Be prepared to justify your answers!
  - What features are important and why?

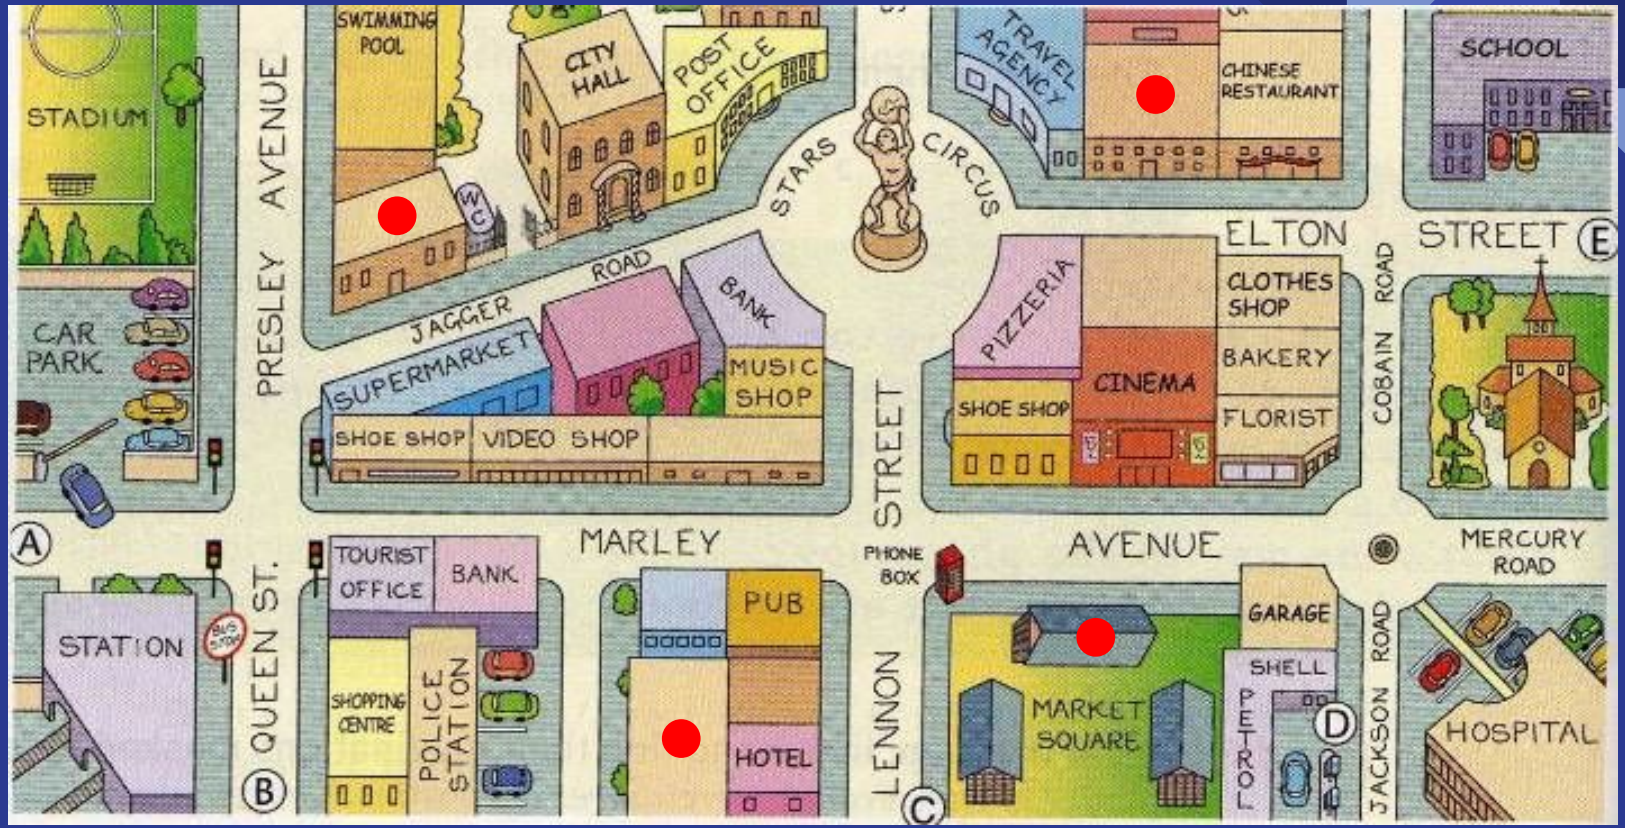
STADIUM
CAR PARK
PRESLEY AVENUE
SWIMMING POOL
CITY HALL
POST OFFICE
WC
STARS
CIRCUS
TRAVEL AGENCY
CHINESE RESTAURANT
SCHOOL
ELTON
STREET
E
JAGGER
ROAD
BANK
SUPERMARKET
MUSIC SHOP
SHOE SHOP
VIDEO SHOP
PIZZERIA
CLOTHES SHOP
BAKERY
CINEMA
SHOE SHOP
FLORIST
COBAIN ROAD
A
STATION
BUS STOP
QUEEN ST.
B
TOURIST OFFICE
BANK
SHOPPING CENTRE
POLICE STATION
MARLEY
STREET
PUB
HOTEL
PHONE BOX
AVENUE
MERCURY ROAD
GARAGE
SHELL
PETROL
MARKET SQUARE
D
LENNON
C
JACKSON ROAD
HOSPITAL

Overview:

- What is clustering?
- Applications of clustering
- Types of clustering
- k-means algorithm
- Recommendation system example

What is clustering?

What is clustering?

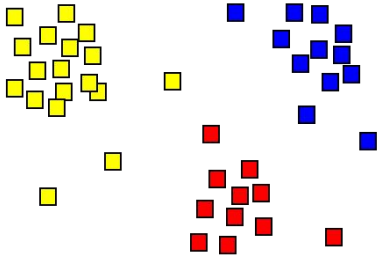

- From Wikipedia:

  “**Cluster analysis** or **clustering** is the task of grouping a set of objects in such a way that objects in the same group (called a **cluster**) are more similar (in some sense) to each other than to those in other groups (clusters).”

- Thinking back to our first lecture on machine learning:

  This is an important type of ***unsupervised learning***. Why?

Applications of clustering

Periodic Table of Elements

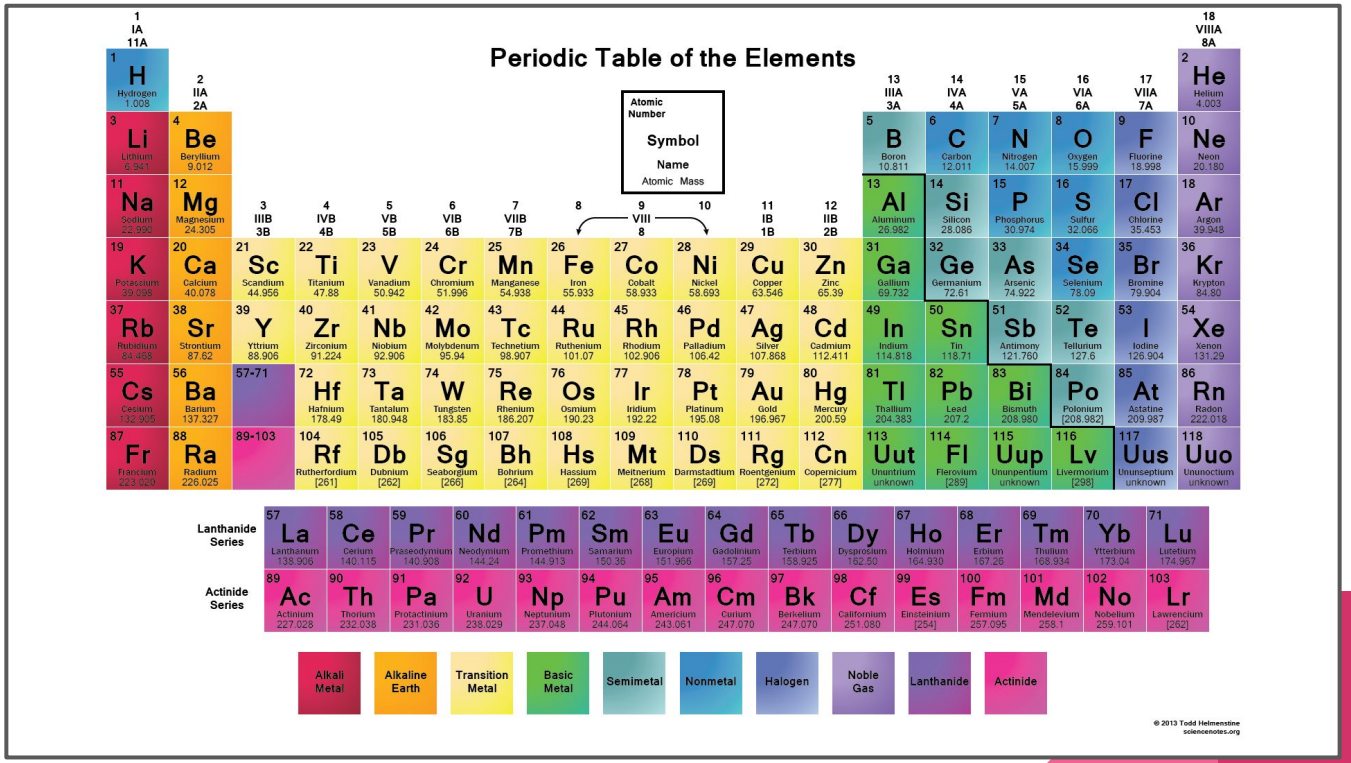

Earthquakes and Seismology

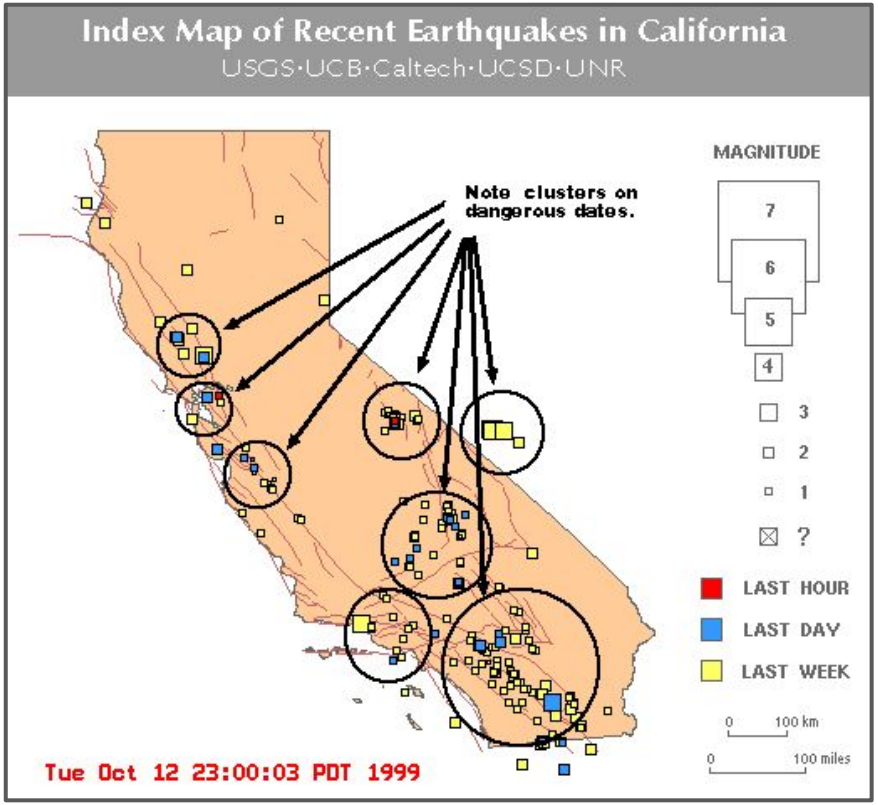

Medical Imaging

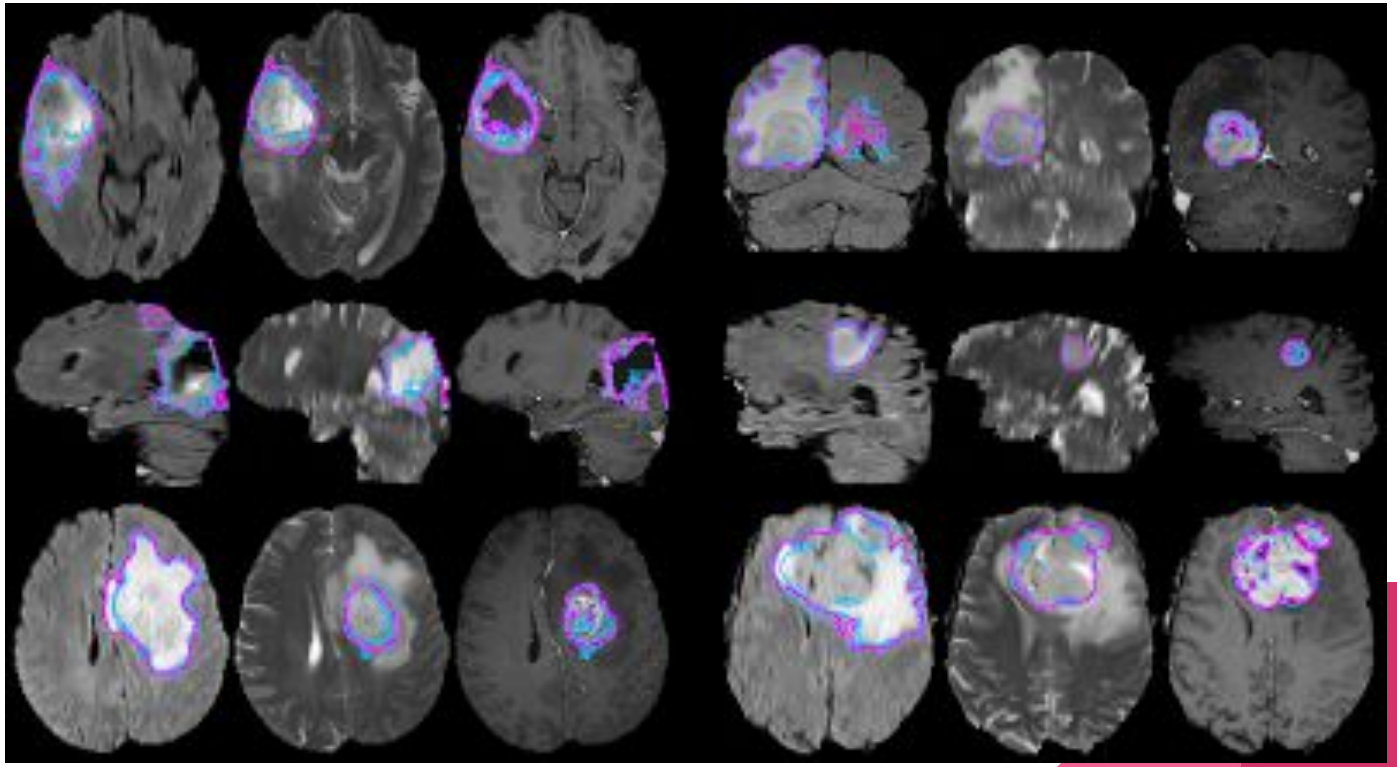

Recommendation Systems: Amazon

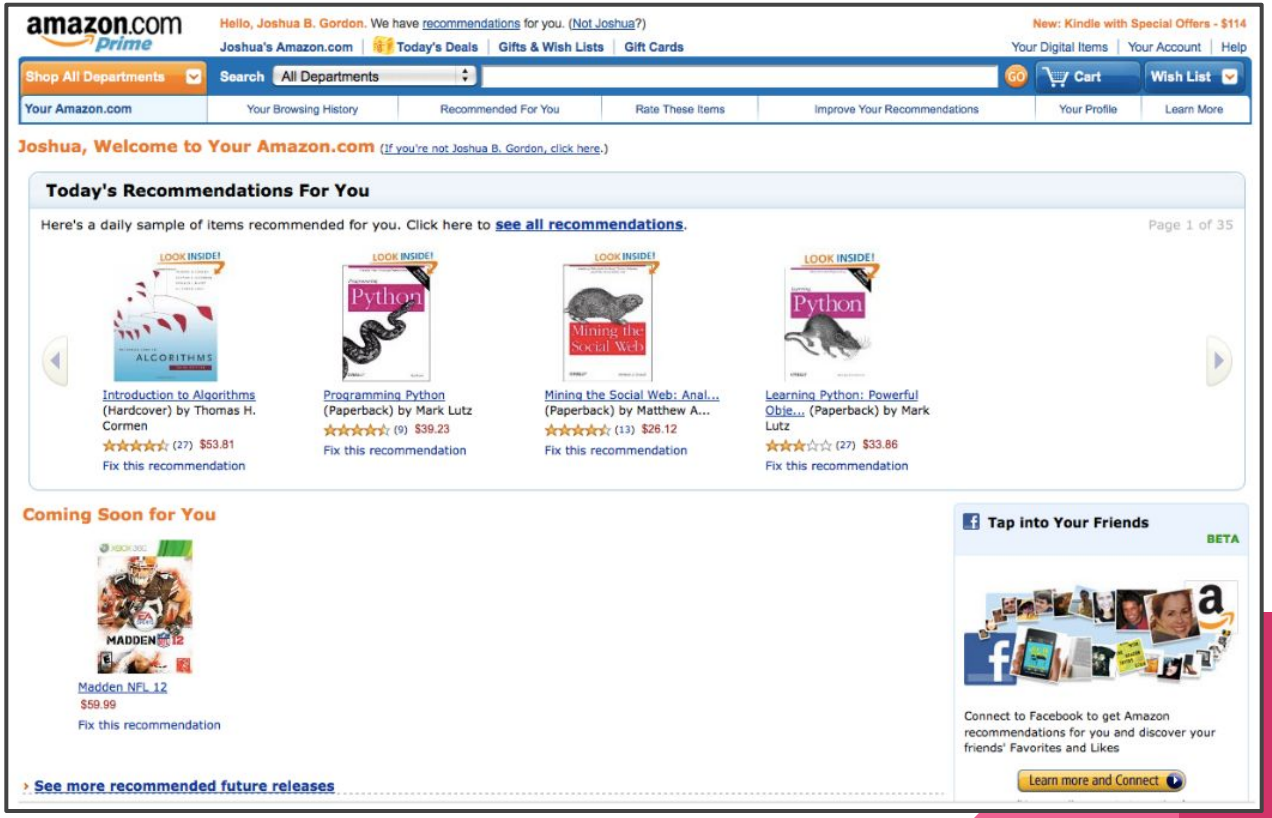

Recommendation Systems: Amazon

- Amazon doesn't know what it's like to read a book, or what you feel like when you read a particular book
- Amazon *does* know that people who bought a certain book also bought other books
- Patterns in the data can used to make recommendations
- If you've built up a long purchase history you'll often see pretty sophisticated recommendations

And even more nefarious purposes...

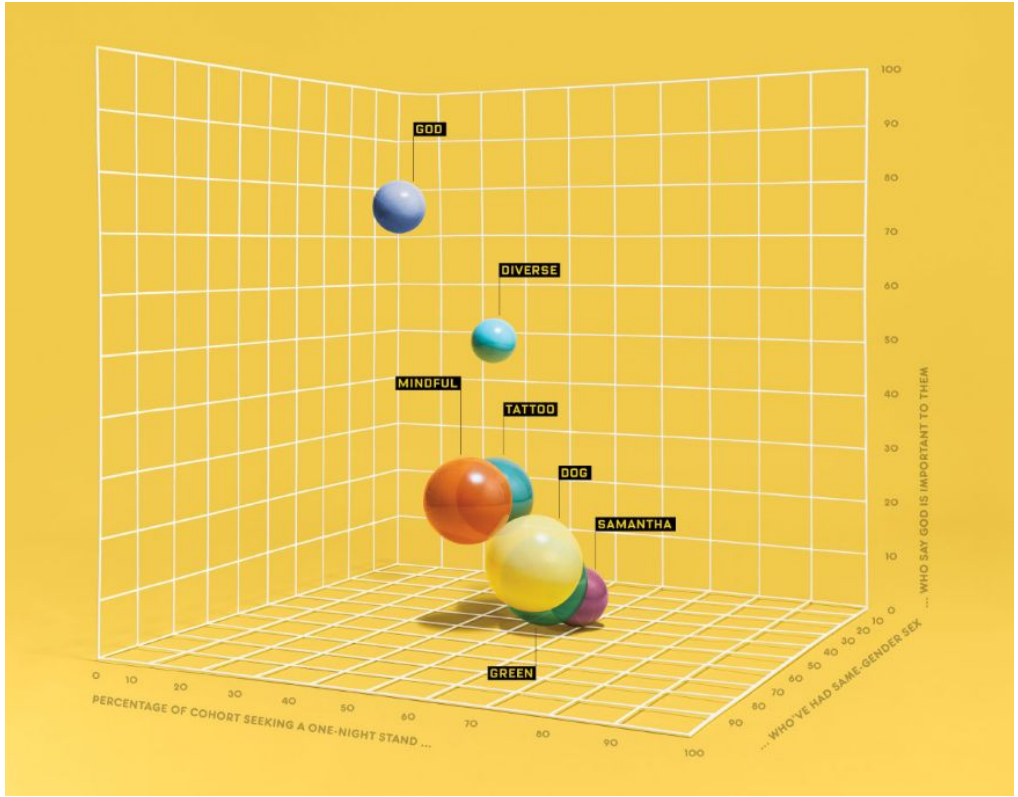

WIRED

KEVIN POULSEN SCIENCE 01.21.14 06:30 AM

HOW A MATH GENIUS HACKED OKCUPID TO FIND TRUE LOVE

https://www.wired.com/2014/01/how-to-hack-okcupid/

Discussion:

Can you come up with other applications of clustering?

Types of clustering

Two main types of clustering:

- **Hard clustering**:

  Each object is in exactly one class.

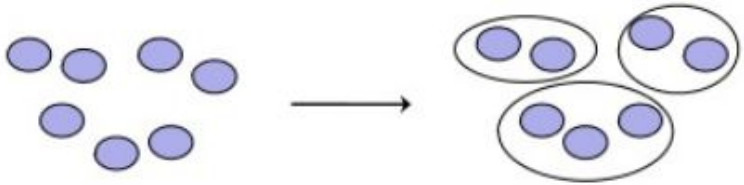

- **Soft clustering**:

  Objects are assigned a degree to which they are in each class.
  - (Think of this as probability of being in a class)

Flat vs. Hierarchical clustering

- Flat = distinct clusters
- Hierarchical = clusters within clusters

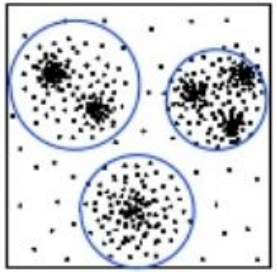

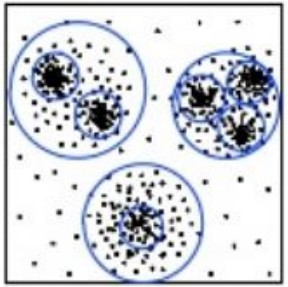

More later on types of clustering *algorithms*.

But first, an example...

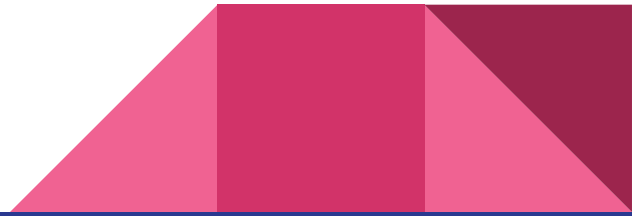

Campgrounds

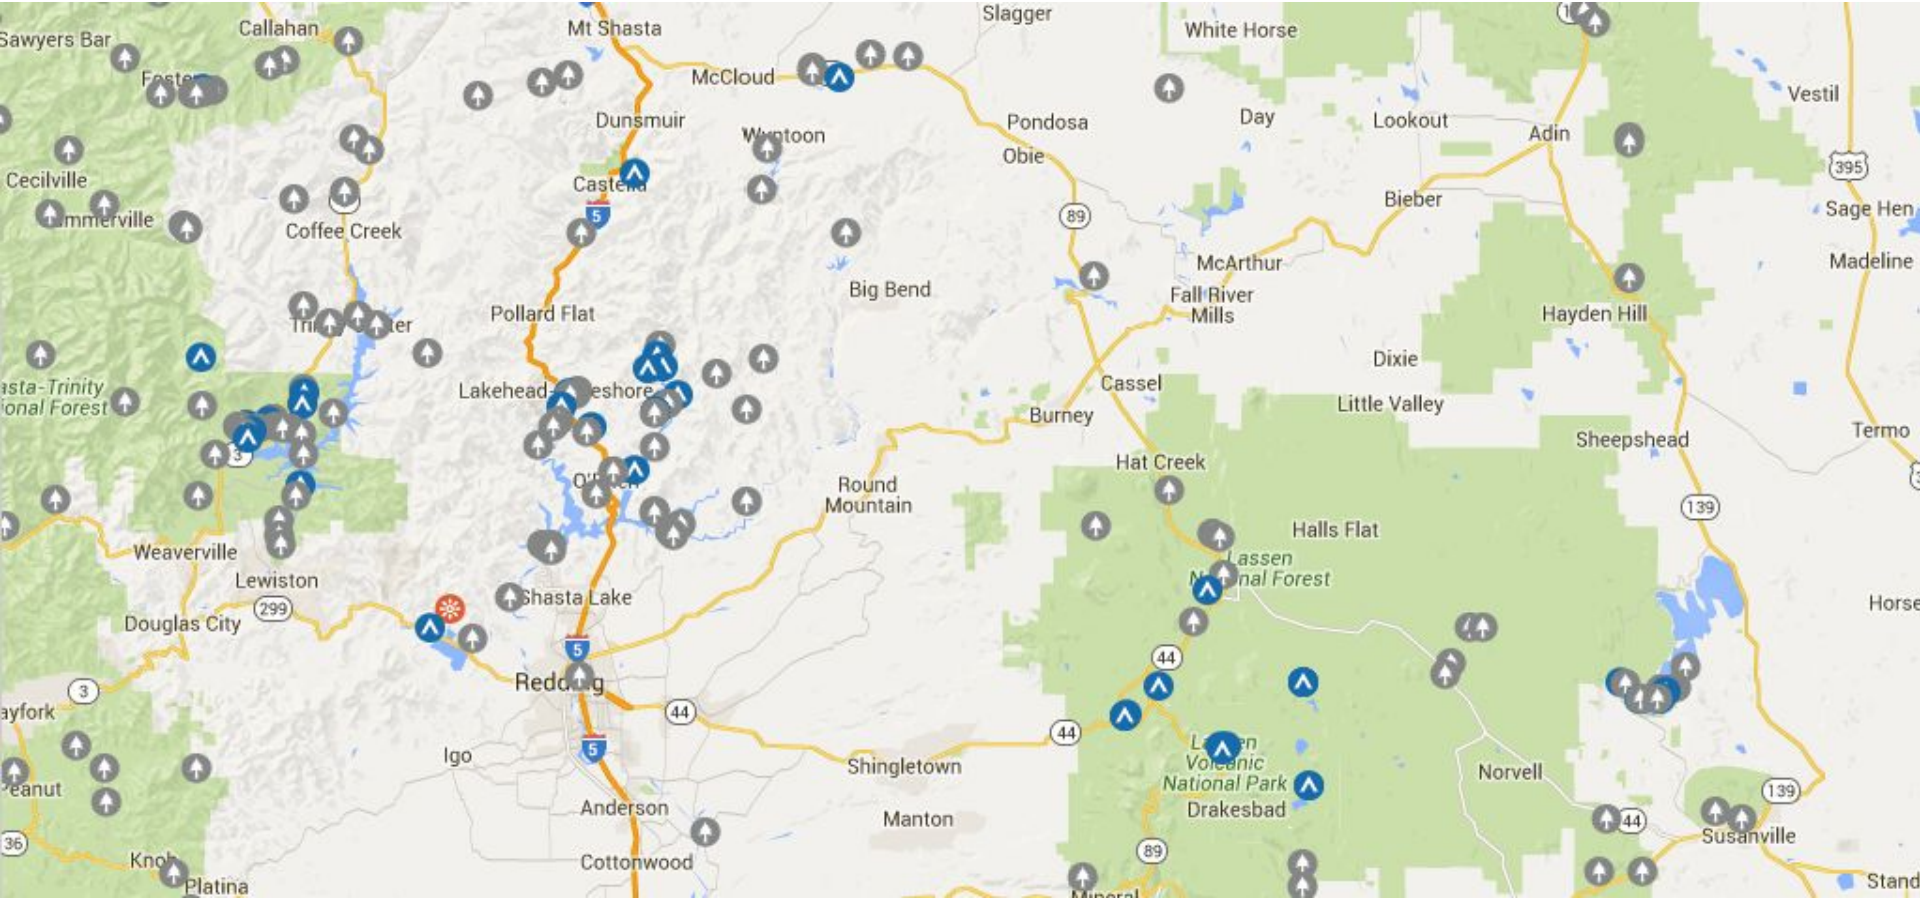

Campgrounds

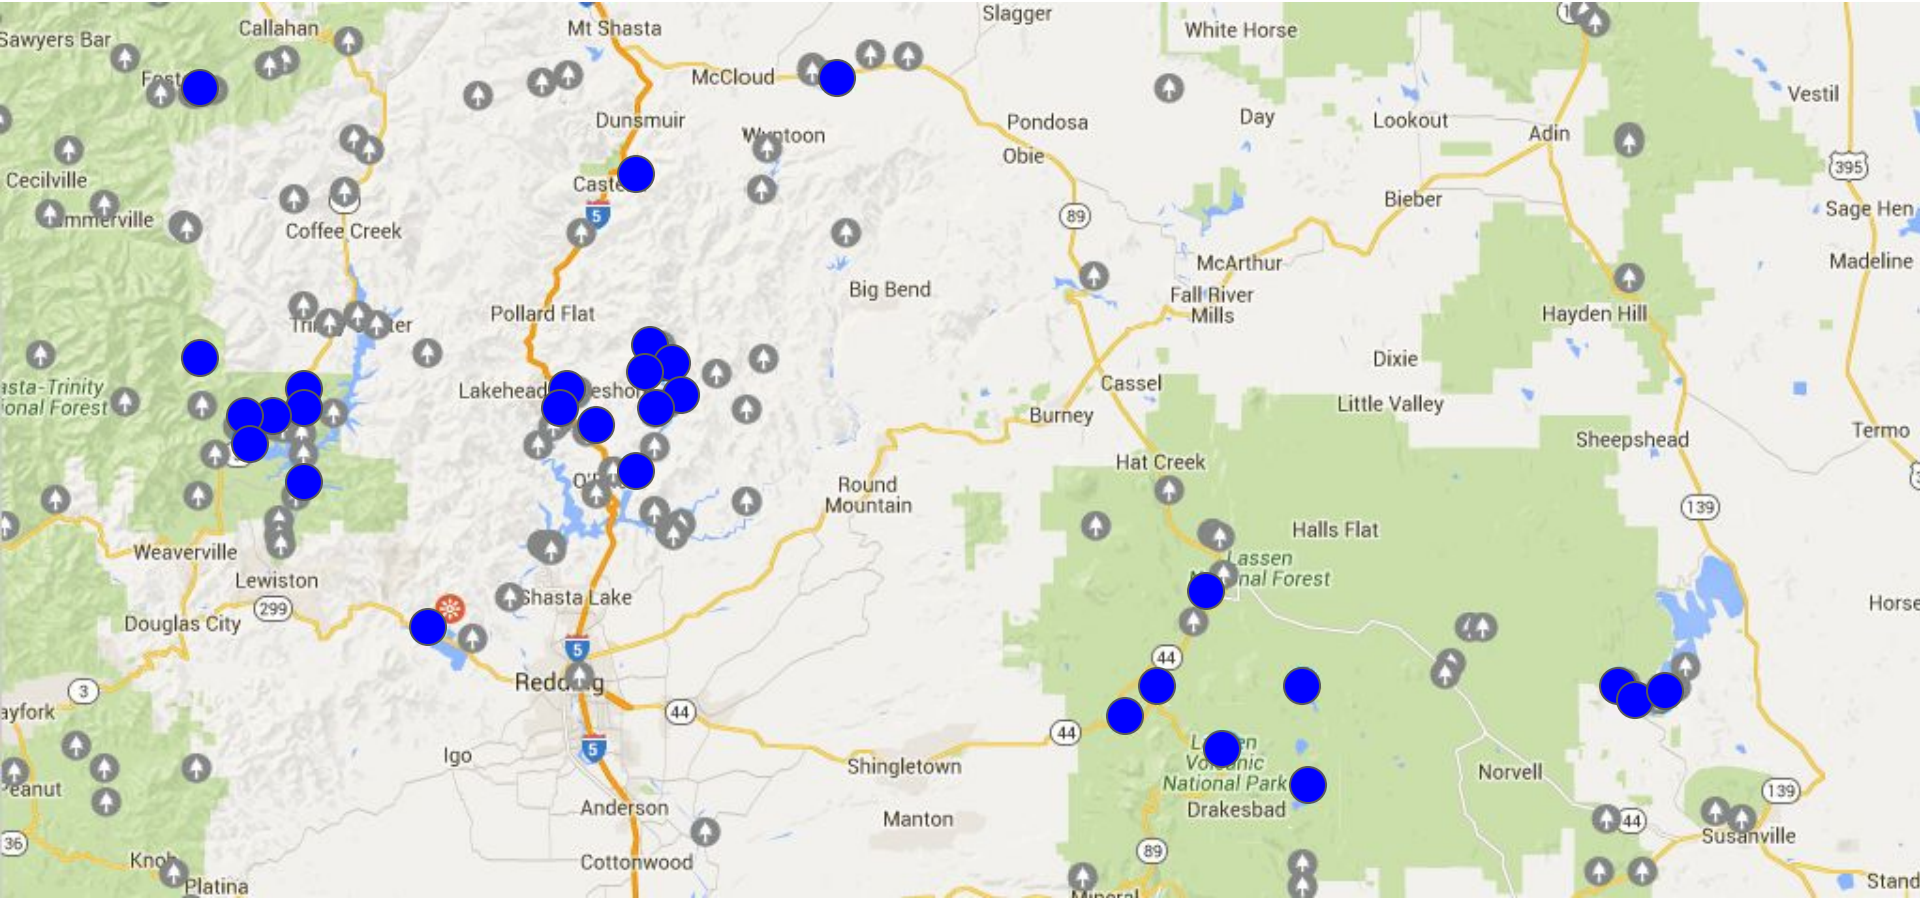

Campgrounds

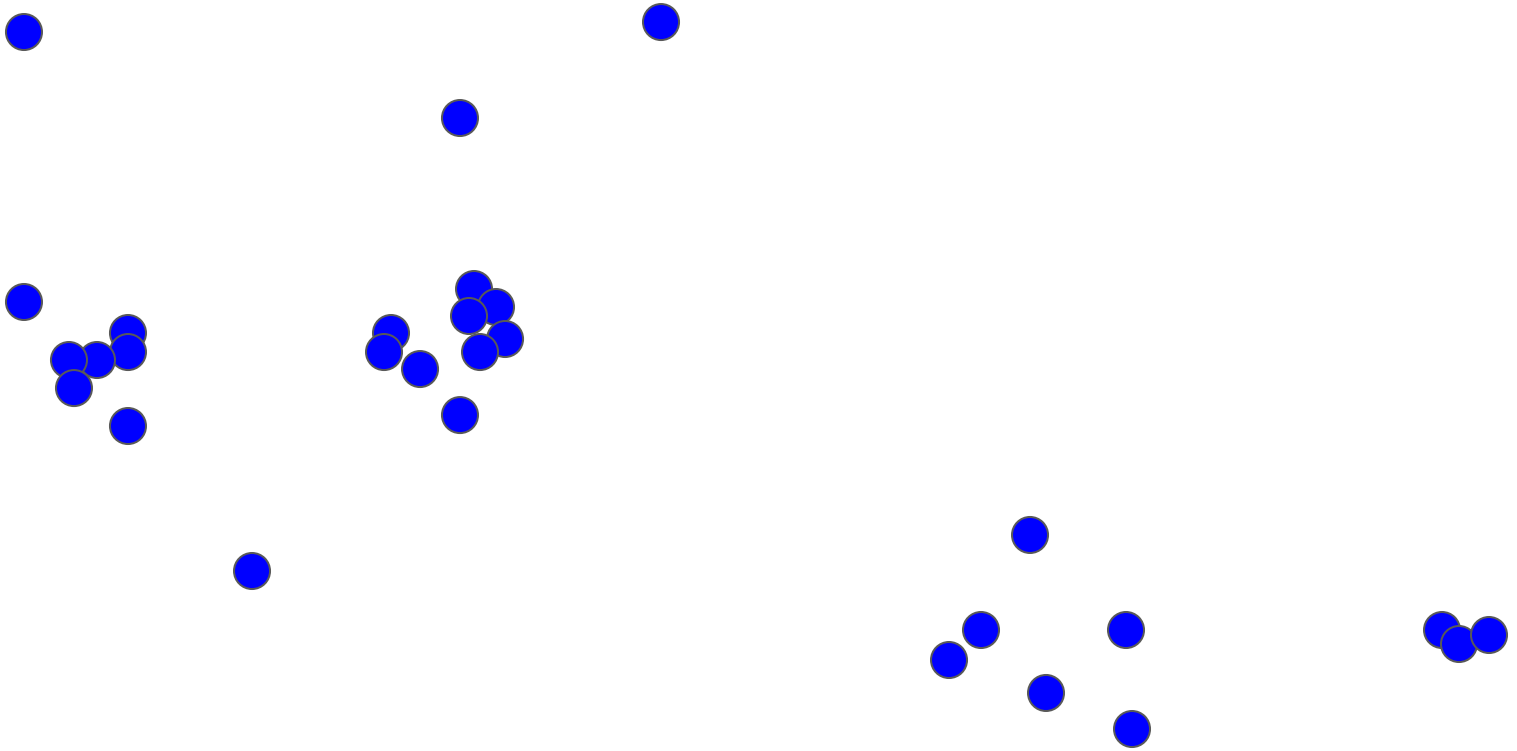

Campgrounds

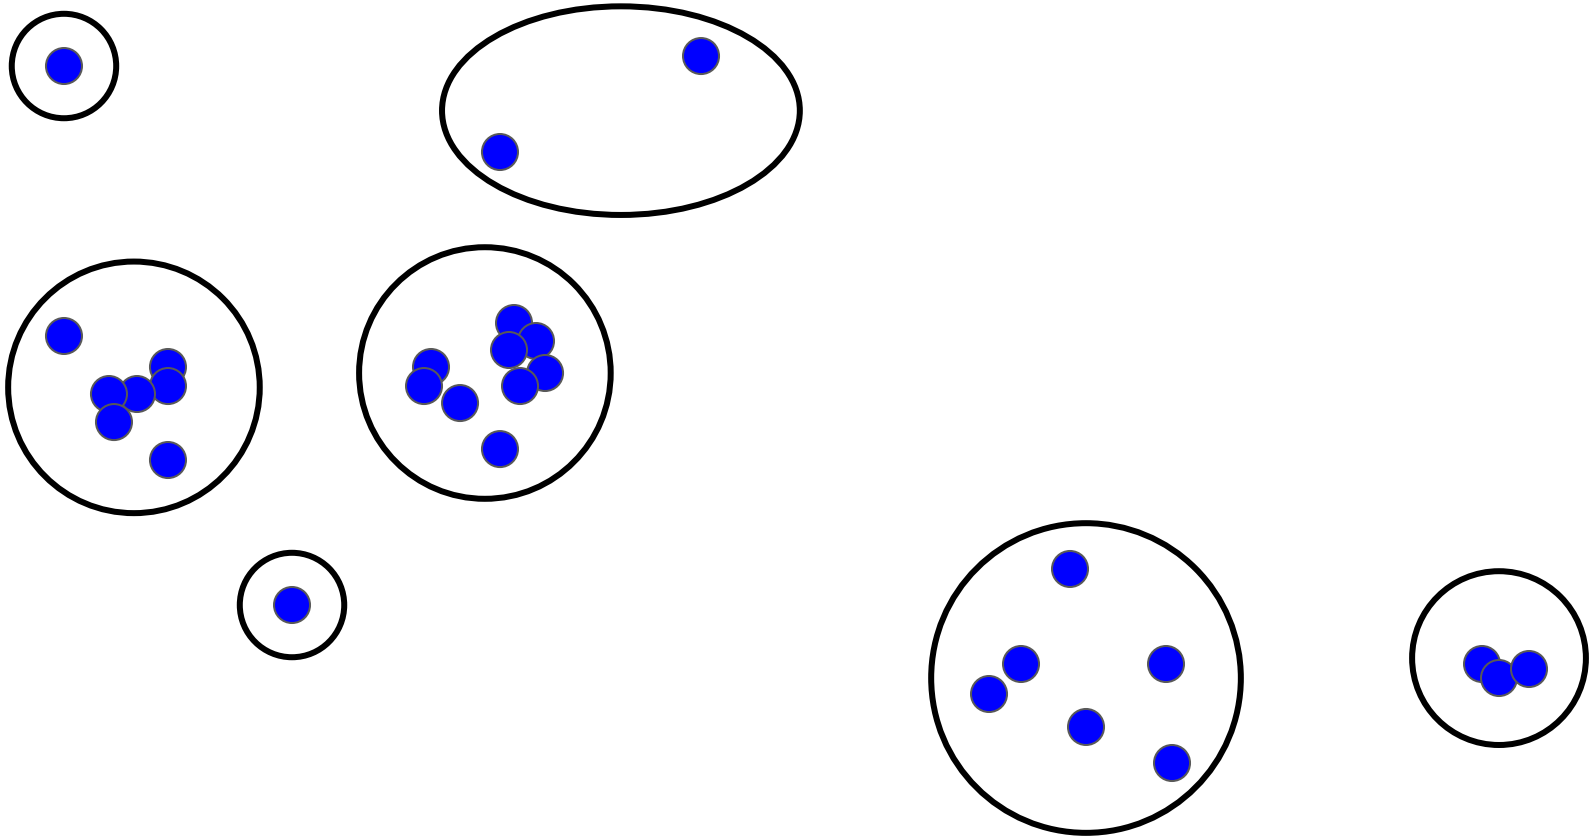

Breakout Session:

Try out the demo at http://mlehman.github.io/kmeans-javascript/

Discuss with your table: how does the algorithm work?

(Note: this demo has a bug! Can’t use more than 6 clusters)

k-means clustering

How can we automatically cluster campgrounds?

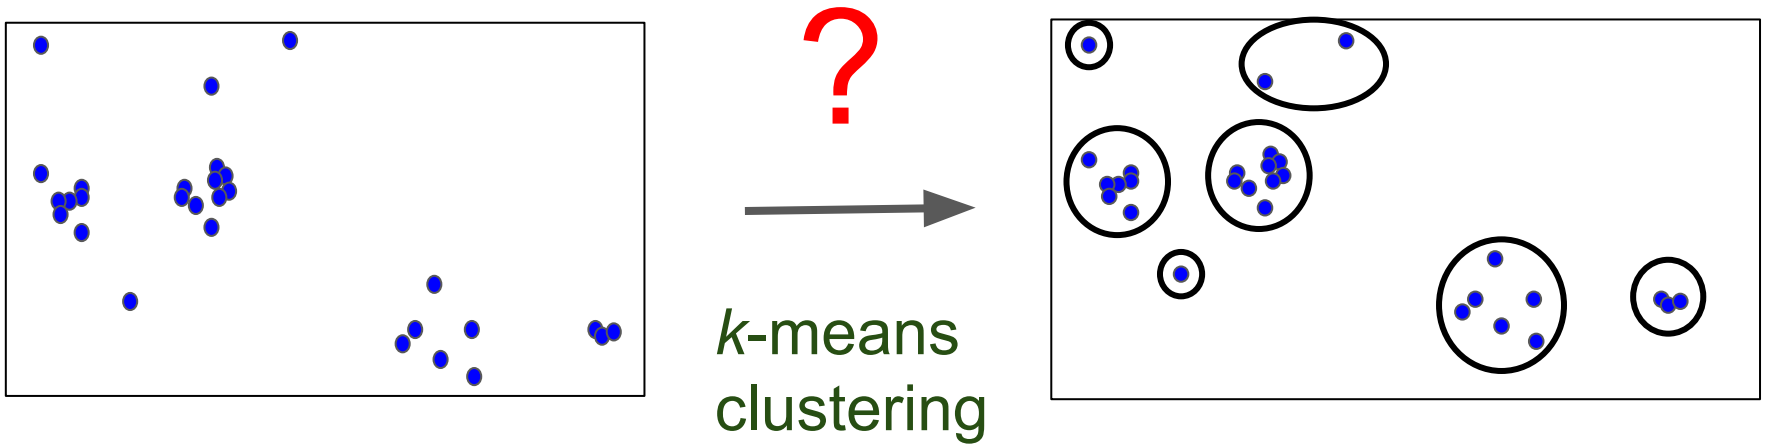

Goal: Assign each of the *n* points to one out of *k* clusters (defined by a cluster centroid).

Find 5 clusters for the 1000 points below!

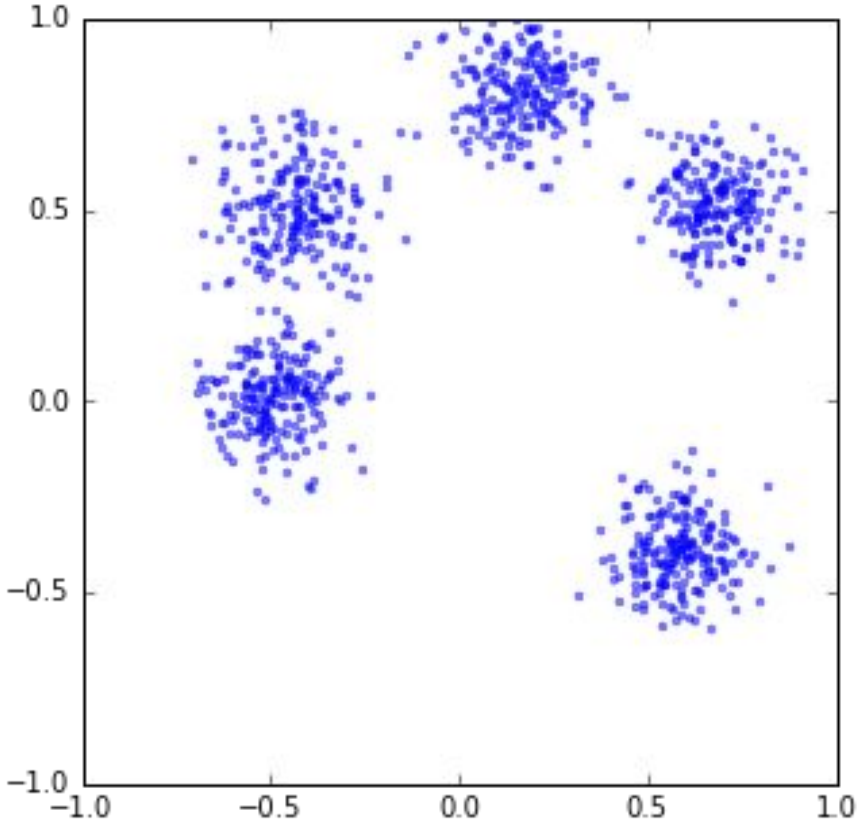

Step 0: Pick k random points as the cluster centroids

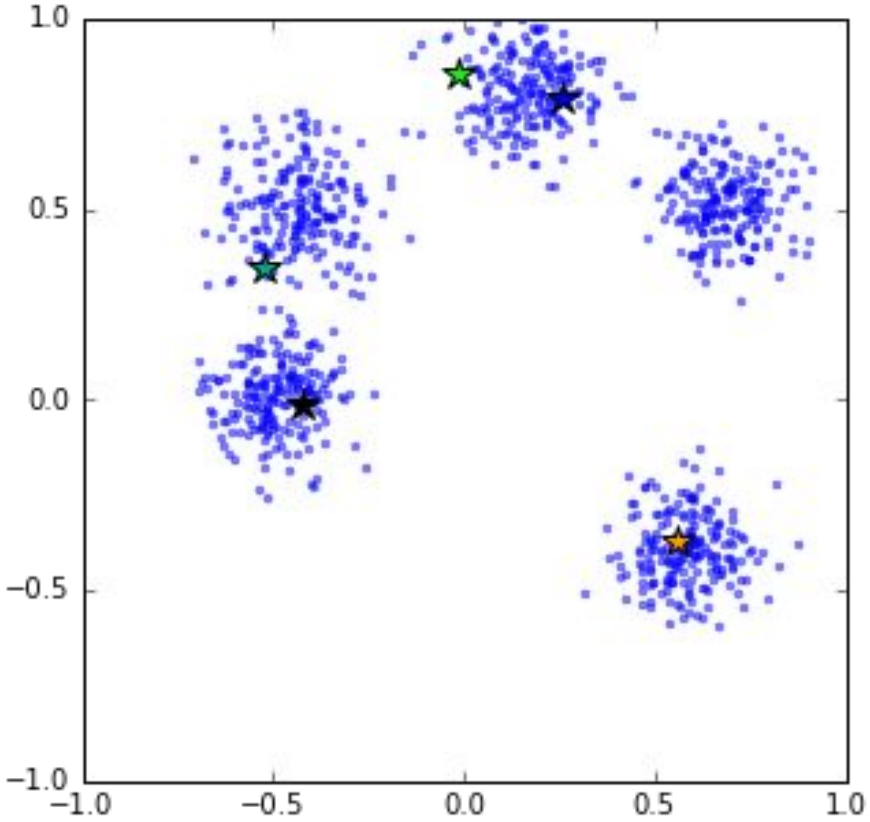

Iteration 1

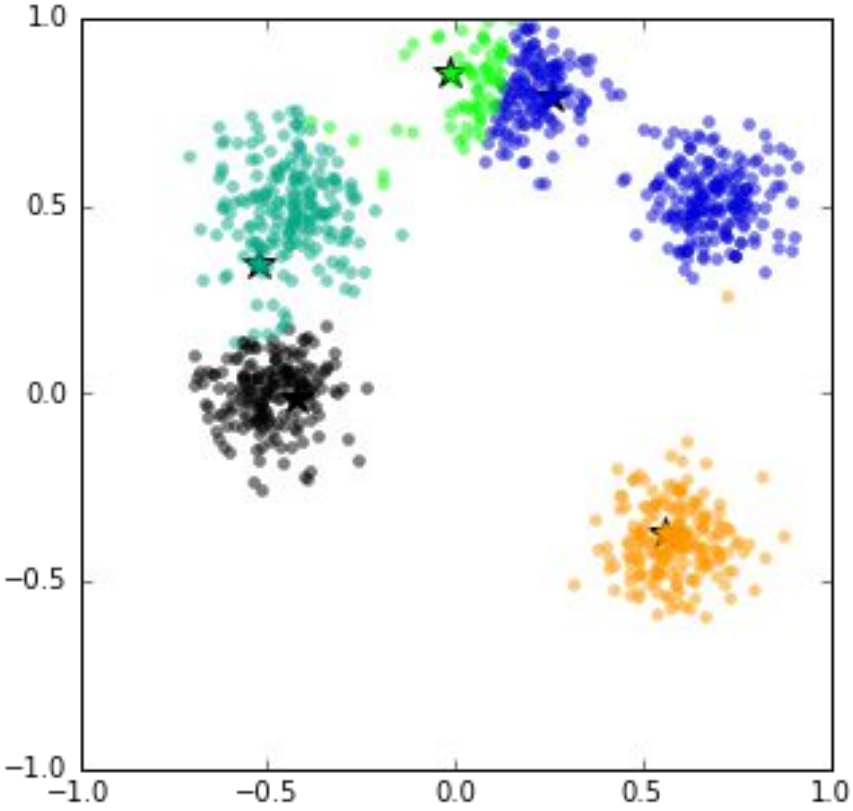

Step 1: Assign each point to closest cluster centroid

Iteration 1 continued

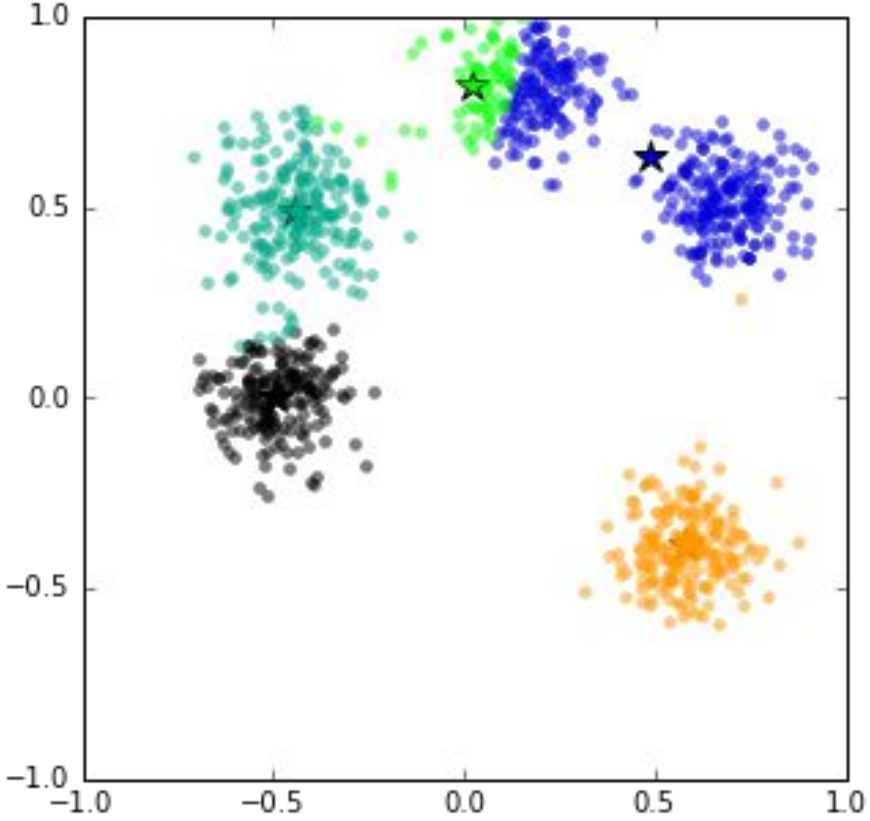

Step 2: Find cluster centroids as center of mass of elements

Iteration 2

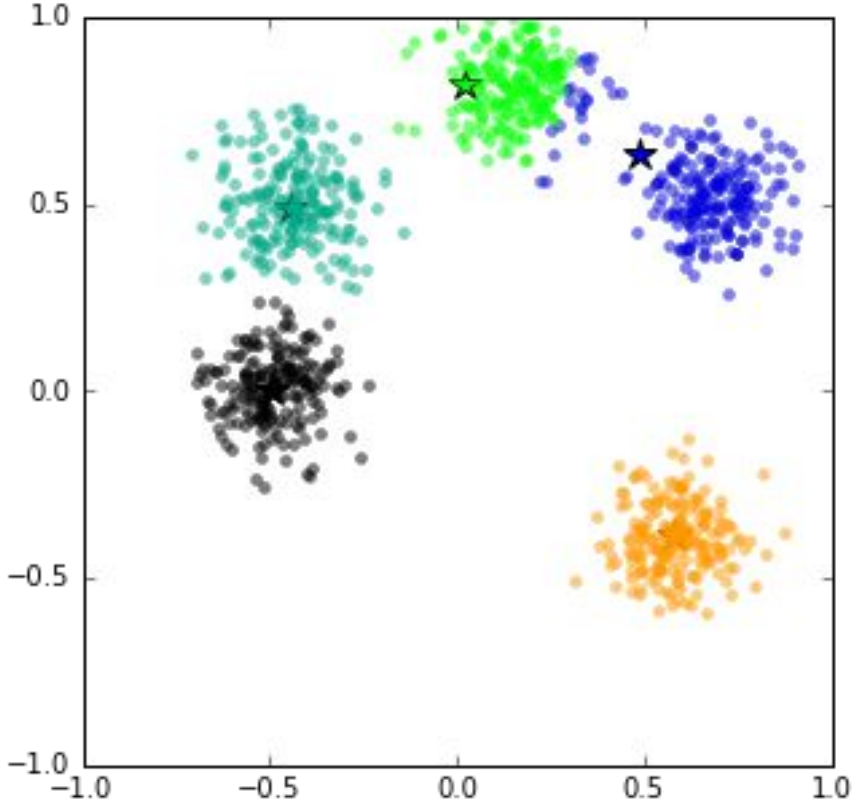

Step 1: Assign each point to closest cluster centroid

Iteration 2 continued

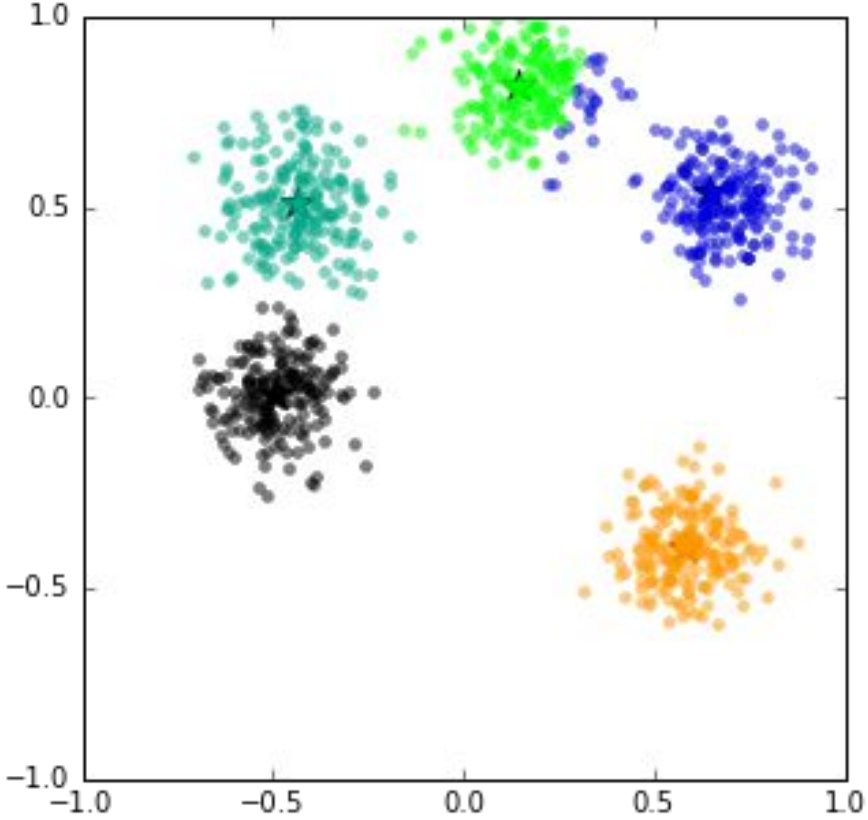

Step 2: Find cluster centroids as center of mass of elements

Iteration 3

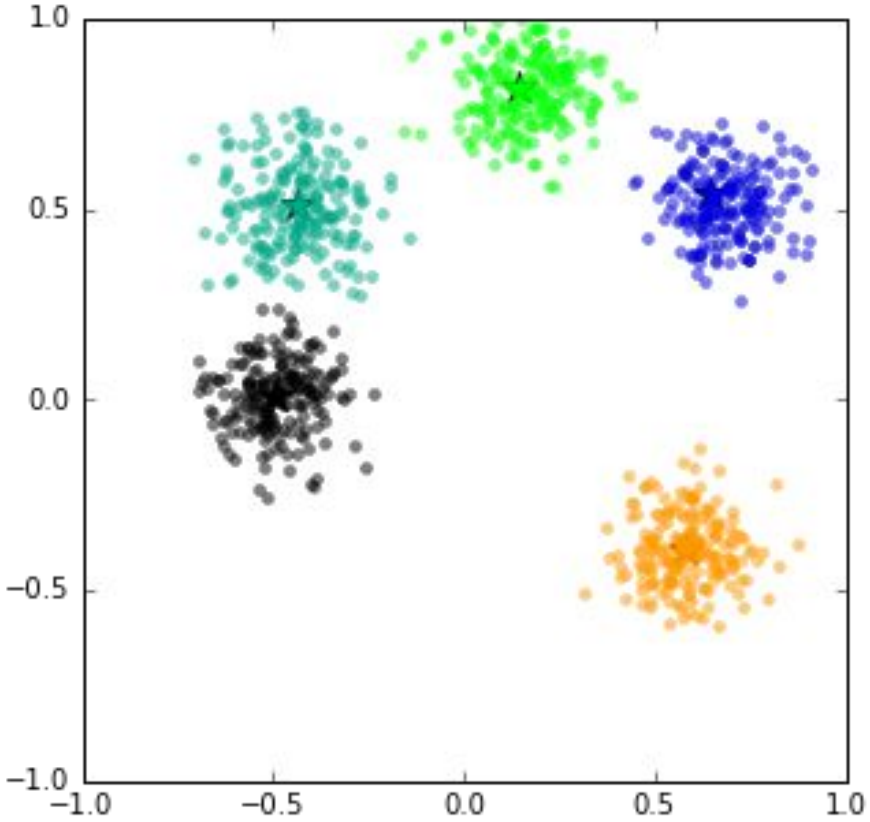

Step 1: Assign each point to closest cluster centroid

Iteration 3 continued

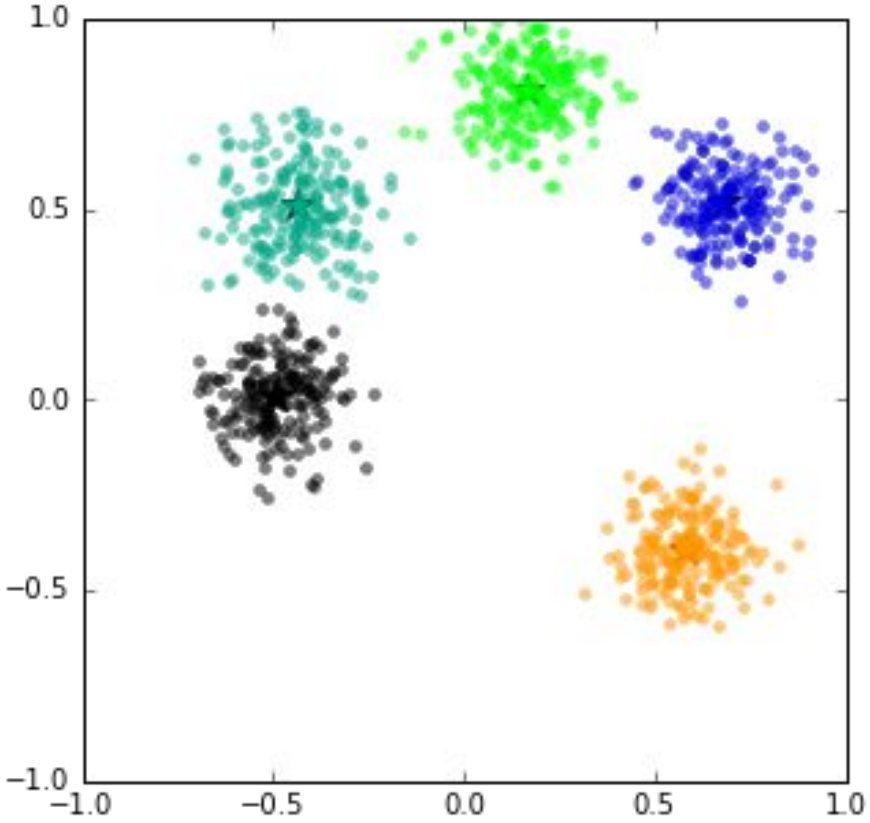

Step 2: Find cluster centroids as center of mass of elements

Iteration 4

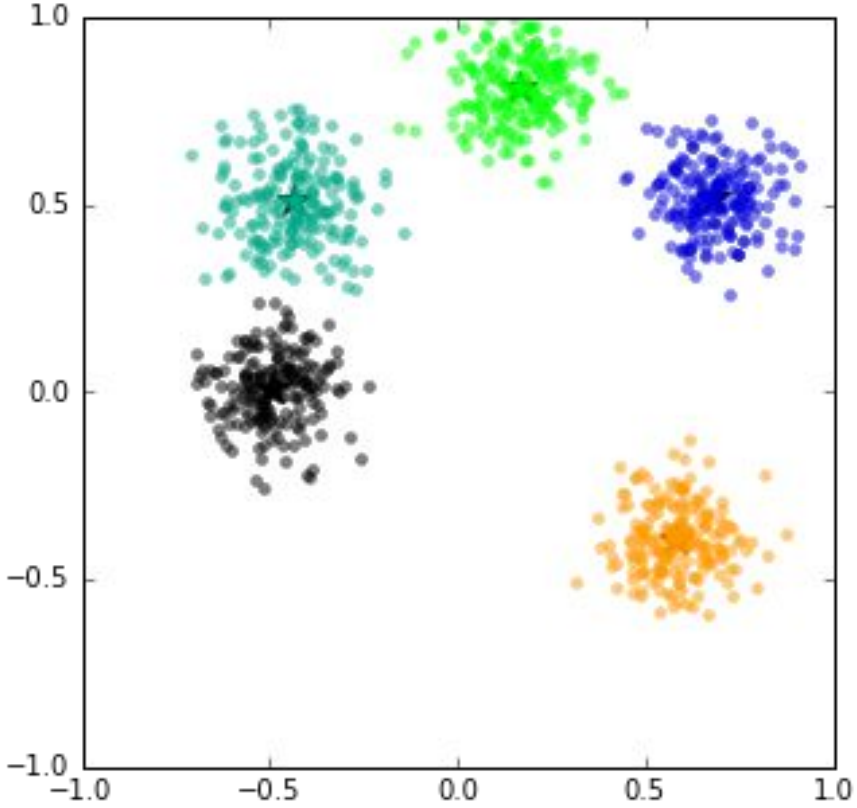

Step 1: Assign each point to closest cluster centroid

No change in cluster assignment.
TERMINATE!

k-means overview

- Choose the value of k
  - This is how many cluster centroids we're finding
- Choose k points in the set
  - These are the initial centroid locations
- For each point not selected, assign to its nearest centroid
  - Now all points have an initial cluster assignment
- Until "happy" do:
  - 1. Recompute centroids of clusters
  - 2. Reassign all points to closest centroid (forms new clusters)

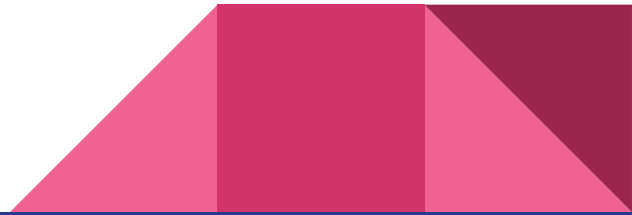

Discussion:

Is this “hard” or “soft” clustering?

Is this “flat” or “hierarchical” clustering?

Discussion:

When would k-means fail?
When would it succeed?

What are some drawbacks?
What are some advantages?

Weaknesses:

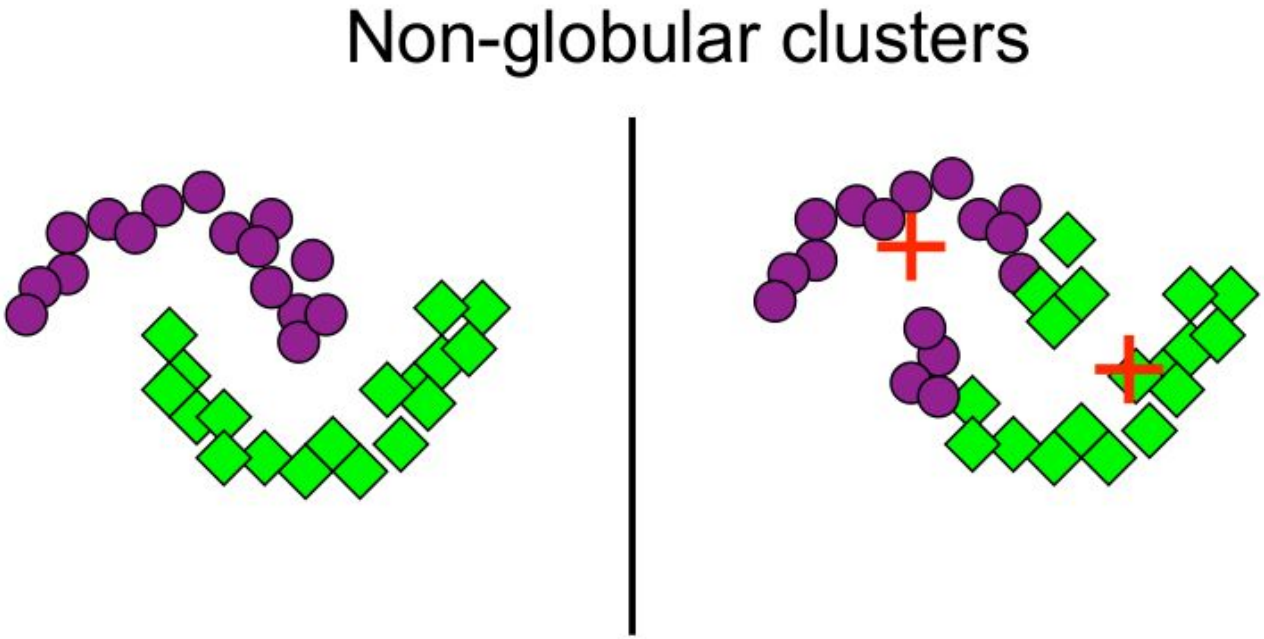

Weaknesses:

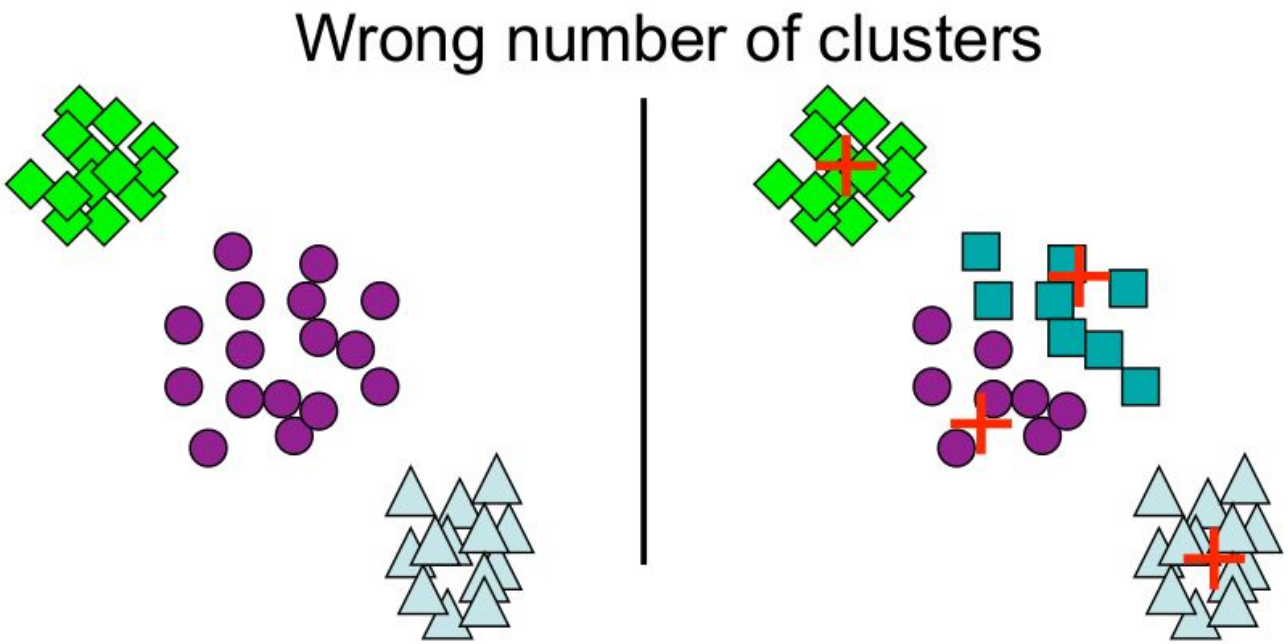

Weaknesses:

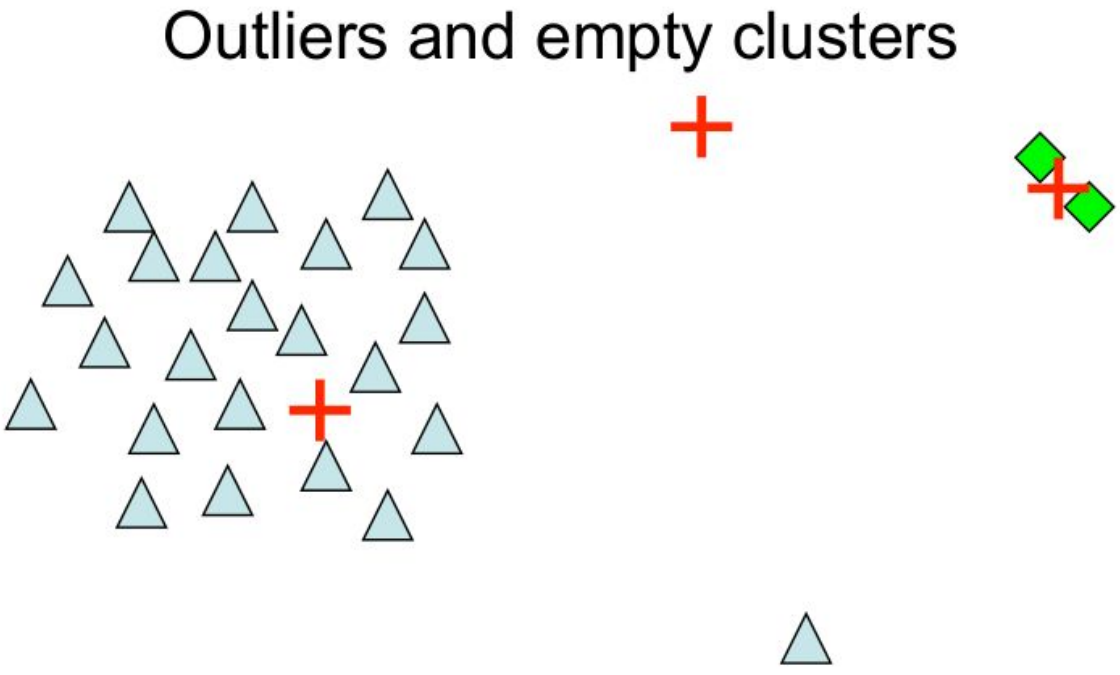

Algorithm design: Distance metrics and features

- How do we determine what elements are “close”?
  - Spatial distance, similarity of campsite attributes,...
- “Features” are attributes that we use to mathematically compute closeness.
  - Height, age, hometown
- Weights for features
  - Are all features equally important or are some more important than others?
  - Are the features on different scales?

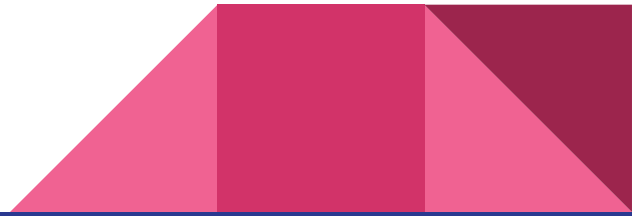

Types of features

- **Categorical features** take a fixed set of values. The values cannot be ordered.
  - For example, the type of phone (android, iphone, windows phone).

- **Ordinal features** take a fixed set of values. The values can be ordered.
  - For example, the ranks of football teams in the PAC-12 conference.

- **Continuous features** can take any real value.
  - For example, the distance of Princeton University from your hometown.

Main types of clustering algorithms

- Connectivity-based (hierarchical) clustering
- Centroid-based clustering
- Distribution-based clustering
- Density-based clustering

Connectivity-based (hierarchical) clustering

- A hierarchy of clusters based on distance:

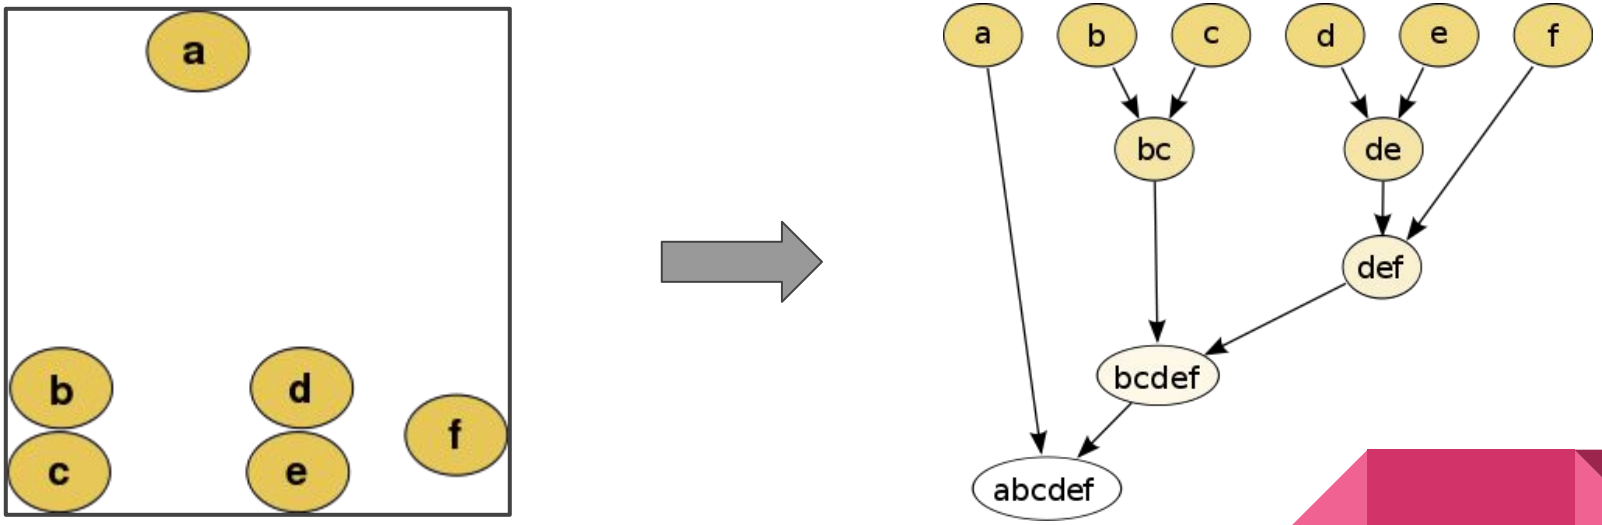

- Typically, form a “dendogram”

Centroid-based clustering

- k-means!

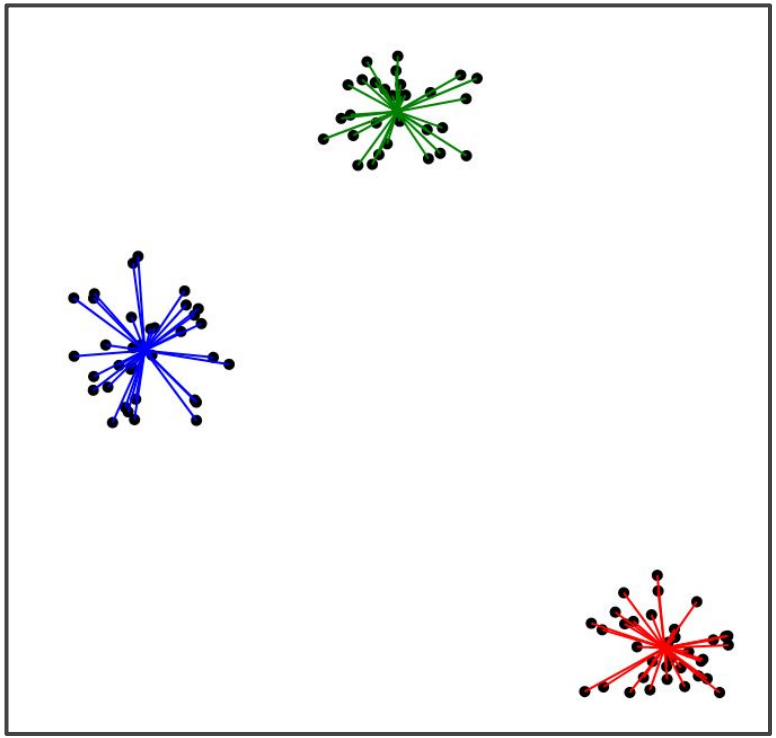

Distribution-based clustering

- Cluster according to how likely it is that points lie in a certain distribution

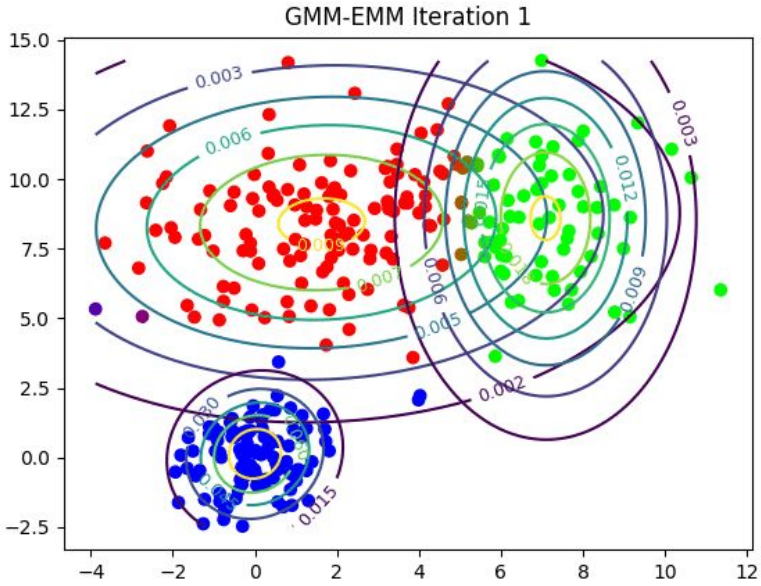

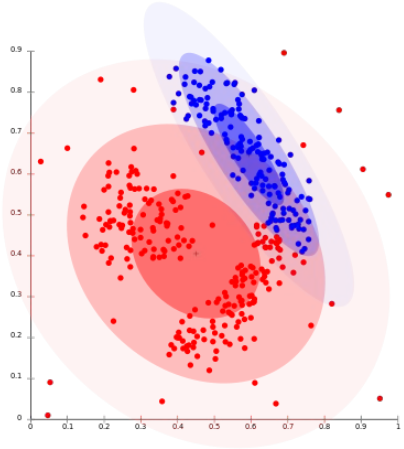

Density-based clustering

- Group together points that are closely spaced, and marks farther-out points as outliers

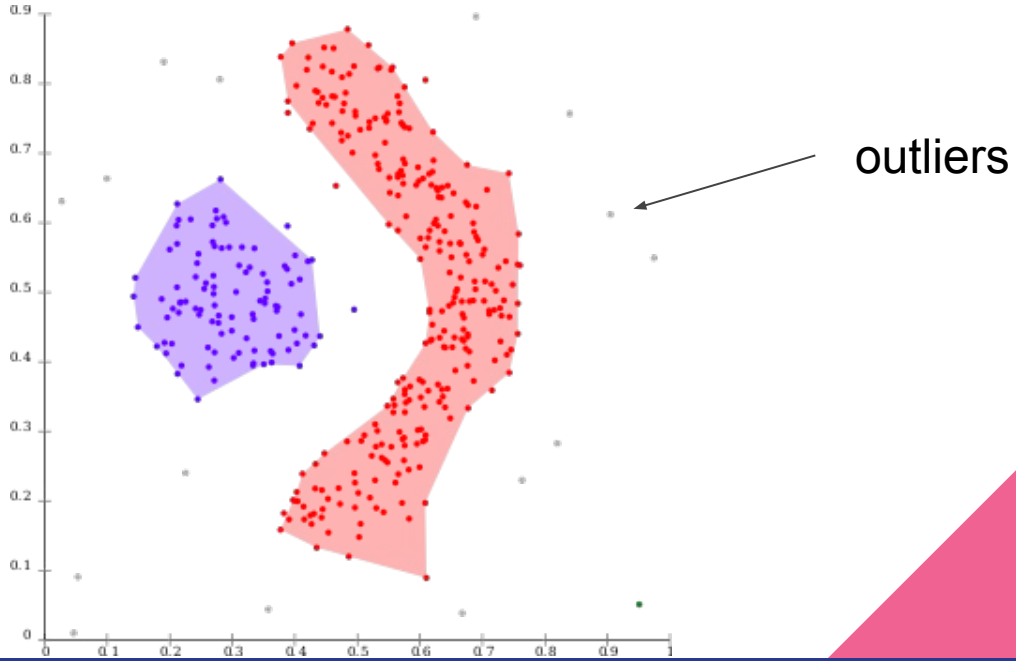

Recommendation system example

Dataset: Movie critics

Critic	Star Wars	Raiders of the Lost Arc	Casablanca	Singin' in the Rain
Sam	****	****	*	**
Sandy	*****	****	**	*
Matt	**	**	****	***
Julia	**	*	***	****
Sarah	*****	?	?	**

- Could an algorithm use this data to recommend movies?
- How would you do it?

Dataset: Movie critics

Critic	Star Wars	Raiders of the Lost Arc	Casablanca	Singin' in the Rain
Sam	****	****	*	**
Sandy	*****	****	**	*
Matt	**	**	****	***
Julia	**	*	***	****
Sarah	*****	?	?	**

- Could an algorithm use this data to recommend movies?
- How would you do it?

Critics with similar taste

- Preference space

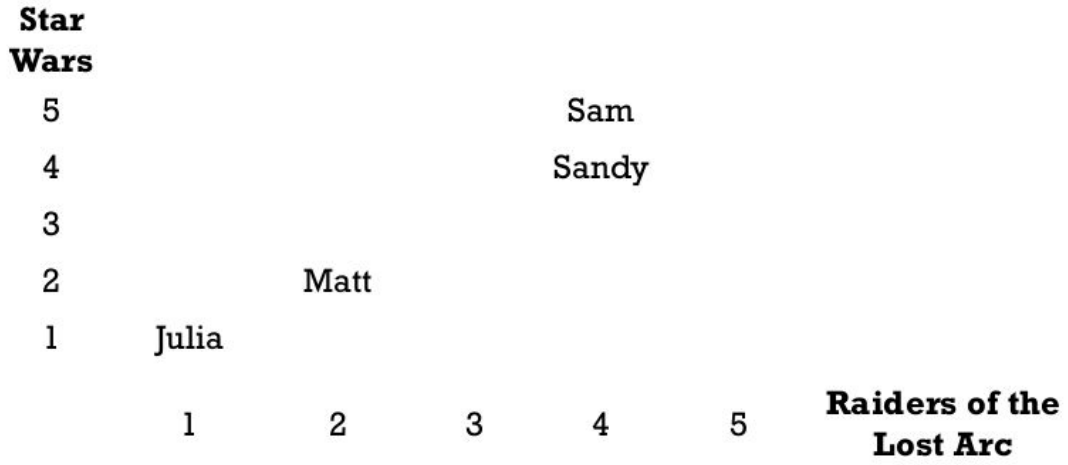

Measuring distance

Star Wars						
5				Sam		
4				Sandy		
3						
2		Matt				
1	Julia					
	1	2	3	4	5	**Raiders**

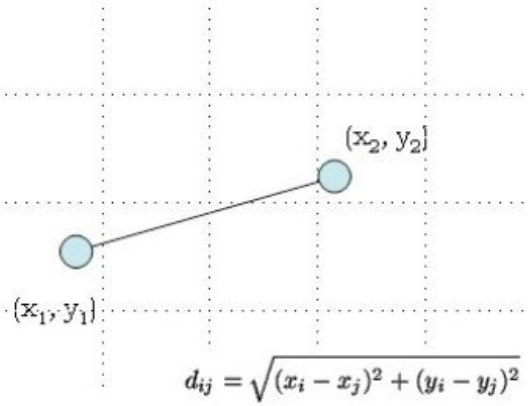

- Measure similarity between points using a measure of distance

Finding critics with similar taste

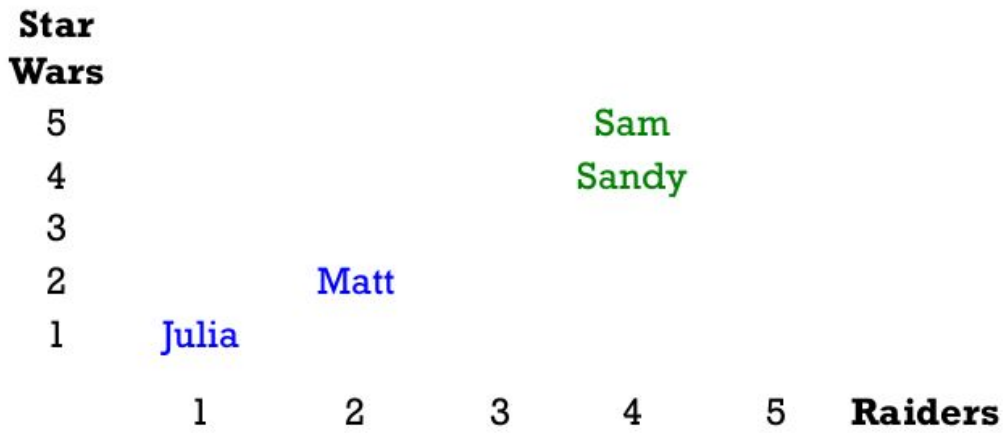

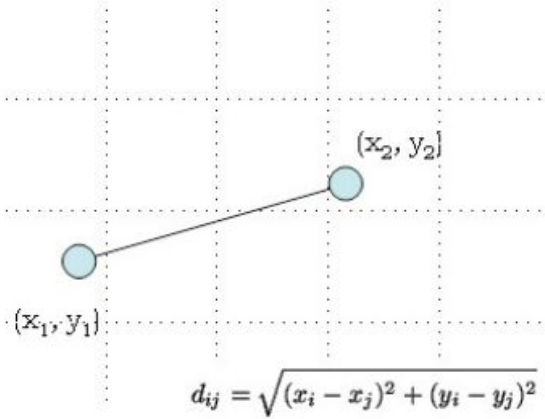

- People who liked Star Wars are close in preference space to those who liked Raiders

Making a recommendation

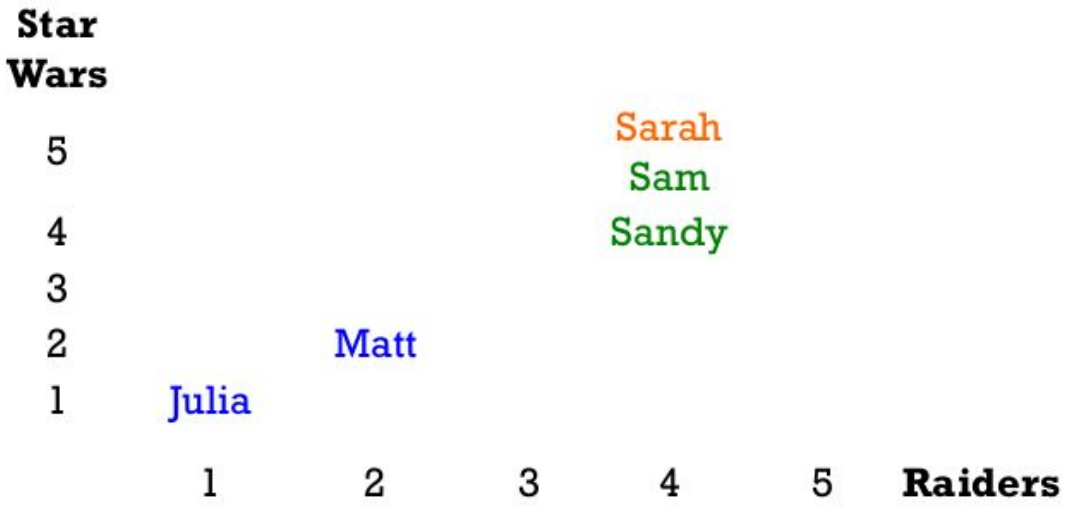

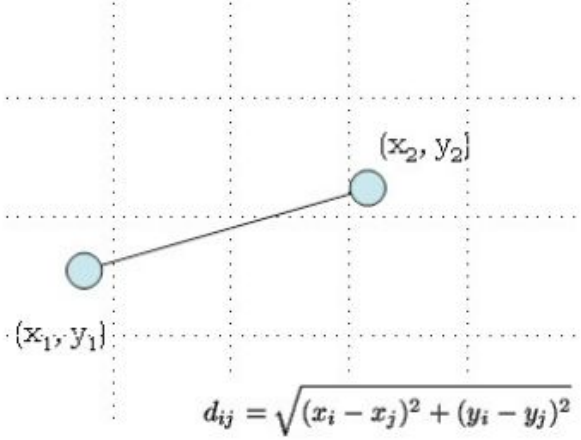

- Sarah hasn't seen Raiders, but gave Star Wars five stars.
- Chances are she'll like Raiders too!

Features

- We used features to compare critics
- Feature: a data attribute used to make a comparison
- Quantify attributes of an object (size, weight, color, shape, density) in a way a computer can understand
- Quality is important

Apples vs. Oranges

- A good feature discriminates between classes
- Think: how well does a feature help us tell two things apart?
- Is mass a good feature? By itself?
- What about in conjunction with another feature like color?

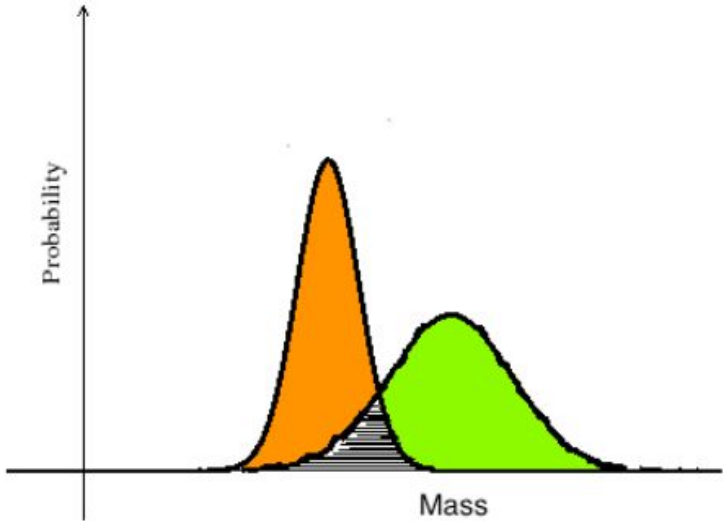

Features to compare movies

Feature	Star Wars	Raiders of the Lost Arc	Casablanca	Singin' in the Rain
…				
…				
…				
…				
…				

- Can you suggest some features to compare movies?

Features to compare movies

Feature	Star Wars	Raiders of the Lost Arc	Casablanca	Singin' in the Rain
Action (1 to 5)	5	4	2	1
Romance (1 to 5)	1	2	4	3
Length (min)	121	115	102	103
Harrison Ford	Y	Y	N	N
Year	1977	1981	1942	1952

- What type of features are these?

Comparing movies in feature space

- Here we can “eyeball” the clusters, but what if we had more features? (i.e. higher dimensional data)?

Exercise time:

Jupyter notebook available at
https://github.com/juliennelachance/ai4all_clustering